Leveraging Additional Resources for Frame-Semantic Role Labeling

Abstract

The high cost of semantic structure annotation is a major obstacle to automating semantic analysis with broad coverage. The fully annotated datasets that exist are often small, hindering the robustness of models trained on them. However, low-resource tasks may benefit from exploiting *partially* annotated data, as well as data with *different* (but related) forms of annotation, for additional training data or features. This paper considers the argument identification and classification subtask of frame-semantic parsing, which to date has relied exclusively upon full-text annotations in the FrameNet resource.[$_S^{NS}$ is this true? I think Dipanjan used exemplars for the latent variable frame ID model, but he hasn't used them at all for arg ID, right?] We augment supervised learning with additional "indirect" training data and features so as to leverage additional resources internal and external to FrameNet (e.g., PropBank). [$_S^{NS}$ result]

1 Introduction

[$_S^{NS}$ sparseness is a challenge for many computational semantics tasks]

Frame-semantic parsing (Das et al., 2014) is a case in point. This is the task of automating the rich linguistic structure analyses of the FrameNet lexicon and corpus (Baker et al., 1998).[1] FrameNet represents kinds of events, scenarios, and relationships with an inventory of **frames** ([$_S^{NS}$ examples]). Each frame is associated with lexical **predicates** (verbs, nouns, adjectives, and adverbs) capable of evoking the scenario, and a set of **roles** (or **frame elements**) called to mind in order to understand

---

[1] http://framenet.icsi.berkeley.edu/

Figure 1: Example sentence from FrameNet full-text annotation. 3 frames and their arguments are shown: BODY_MOVEMENT is evoked by *crane*, OBSERVABLE_BODY_PARTS by *necks*, and PUNCTUAL_PERCEPTION by *glimpse*. (Further, *harbor* is annotated as evoking the LOCALE_BY_USE frame and doubles as its sole argument.) Horizontal lines representing argument spans are labeled with role names.

the scenario. These roles may be implicit, but are frequently realized linguistically in the same sentence as the predicate. Given a sentence, frame-semantic parsing is the task of mapping tokens in the sentences to evoked frames, and for each evoked frame, finding and labeling its **argument** phrases with roles. An example appears in figure 1; it will be explained in detail in §2.2.

FrameNet 1.5 defines a structured hierarchy of over 1,000 frames associated with [$_S^{NS}$ #] English lexical predicates, and also provides annotations for [$_S^{NS}$ # targets annotated total] attestations of these frames/predicates in corpora, annotated in context with their arguments. In FrameNet 1.5, a rather small number of sentences—[$_S^{NS}$ #], comprising [$_S^{NS}$ #] words—are provided with **full-text** annotations, i.e. the sentence has been analyzed for all available frames. But a full [$_S^{NS}$ #]% of sentences in FrameNet—the lexicographic **exemplars**—are annotated for only one frame per sentence, and have thus far not been exploited successfully for frame-semantic parsing. Here, we seek to leverage these exemplar sentences as well as the (type-level) hierarchical structure of the FrameNet lexicon.

In this paper, we address the **argument identification** subtask of finding and labeling arguments given a predicate in context and the frame it evokes.

Figure 2: Ideal PropBank annotations for verbs. Though PB uses lexical frames rather than deep frames, there are clear similarities to the FrameNet annotations in figure 1.

This is a form of semantic role labeling (SRL), a task introduced by **?** using a much earlier version of FrameNet. Notably, another resource, **PropBank** (Kingsbury and Palmer, 2002), has been widely used for SRL (Palmer et al., 2010). PropBank annotations capture shallower lexical frames and arguments; additionally, PropBank provides [$^{NS}_S$ millions?] of words of fully annotated English sentences (annotation is much less expensive, but also potentially less valuable, because of the shallower representation). Despite a number of differences in the representations and annotation conventions, for many predicates FrameNet and PropBank are really quite similar: figure 1 and figure 2 show this for the verb *crane*. To get the best of both worlds, we aim to tap into PropBank's vast resources as indirect token-level supervision for FrameNet-style analysis. We hypothesize that PropBank analyses can serve as a weak signal for the FrameNet SRL task, either by heuristically transforming PropBank annotations into FrameNet annotations to augment the training data, or by preprocessing sentences with a PropBank SRL system to obtain new features for FrameNet argument identification.

Our experiments expand the *training data* and/or the *feature space* of supervised argument identification in order to integrate evidence from all of these sources into SEMAFOR (Das et al., 2014), the leading open-source frame-semantic parser for English.[2] The results show that some of these sources of evidence succeed at boosting argument identification performance.[$^{NS}_S$ SOTA (without constraints)?]

2 Resources

[$^{NS}_S$ incl. data analysis of differences from full-text]

[$^{NS}_S$ w/in each mention: genre/overall vocabulary; coverage and distributions of predicates, FE labels; oracle coverage of FT test]

---

[2] http://www.ark.cs.cmu.edu/SEMAFOR/

2.1 The FrameNet Lexicon

Each frame in the Berkeley FrameNet lexicon is intended to represent a gestalt scene. The frame definition includes: a descriptive name; a set of **core roles** representing participants and props that are crucial to understanding the scene; a set of **non-core roles** such as circumstantial information (time, place, manner, purpose, etc.); an English textual description of the scene and how its roles relate to one another; and a set of English predicates that can evoke the scene. For example, the BODY_MOVEMENT frame has **Agent** and **Body_part** as its core roles; the frame description states, "This frame contains words for motions or actions an **Agent** performs using some part of his/her body." Lexical entries in this frame include verbs such as bend, blink, crane, and curtsy, plus the noun use of curtsy.

The frame lexicon is organized as a network, with several kinds of **frame-to-frame relations** linking pairs of frames and (subsets of) their arguments (**?**). Among these kinds of frame-to-frame relations are:

- **Inheritance:** E.g., PUNCTUAL_PERCEPTION (e.g., glimpse.v) inherits from PERCEPTION_EXPERIENCE (e.g., see.v), which inherits from PERCEPTION. Other frames inheriting from PERCEPTION include SENSATION (e.g., sight.n) and BECOMING_AWARE (e.g., notice.v). Crucially, roles in inheriting (conceptually more specific) frames are mapped where they correspond to a role in the inherited frame: so PUNCTUAL_PERCEPTION.**Perceiver** links to PERCEPTION_EXPERIENCE.**Perceiver_passive**, which links to PERCEPTION.**Perceiver**, which links to SENSATION.**Perceiver_passive** and BECOMING_AWARE.**Cognizer**.

- **Subframe:** This indicates a subevent within a complex event. E.g., the CRIMINAL_PROCESS frame groups together subframes ARREST, ARRAIGNMENT, TRIAL, and SENTENCING. CRIMINAL_PROCESS.**Defendant**, for instance, is mapped to ARREST.**Suspect**, ARRAIGNMENT.**Defendant**, TRIAL.**Defendant**, and SENTENCING.**Convict**. Other salient participants in the complex event (such as the crime for which someone is arrested, tried, etc.) are similarly mapped via Subframe relations. This permits an inference that a person tried for a crime likely has been or will be arrested, arraigned, and sentenced for that crime.

In §3.4, we experiment with features shared between related roles of related frames in order to capture statistical generalizations about the kinds of arguments seen in those roles and how they relate syntactically to the predicate.

2.2 Full-text Annotations

Contemporary frame-semantic parsers are trained and evaluated on the **full-text (FT)**[3] portion of the FrameNet corpus. This consists of documents for which annotators made an effort to assign frames and arguments to as many words as possible. Figure 1 gives an example sentence from the FT portion of the corpus. It has 4 frame annotations. BODY_MOVEMENT, as described in the previous section, is evoked by *crane*, and 3 of its roles are filled by overt arguments: the 2 core roles (**Agent**, **Body_part**) happen to be filled by noun phrases (*passengers*, *their necks*), while the non-core role **Purpose** is filled by a prepositional phrase adjunct (*for dizzying glimpses of the harbor*). The frame defines 19 additional non-core roles, none of which have an argument in the example. In frame semantics, non-core roles are considered to be *conceptually* optional; core roles may or may not be *syntactically* optional, but if not locally specified they are expected to be available from context, or else implicit. For example, PUNCTUAL_PERCEPTION— evoked in this case by *glimpses*—is annotated as missing 2 of its core roles. A human listener would resolve the identity of the **Perceiver** from the wider context and the **Body_part** from world knowledge.

In some cases, FT annotation involved creating a new frame or adding a new lexical unit to an existing frame. In other cases, words that in principle should be considered to evoke a frame were left unannotated because they did not match any existing lexical units. This was the case for *passengers* and *dizzying* in figure 1.

[$^{\text{NS}}_{\text{S}}$ Genres, sizes]

[$^{\text{NS}}_{\text{S}}$ Annotation density: (proportion of tokens evoking a frame. breakdown by POS?)]

2.3 Exemplars

[$^{\text{NS}}_{\text{S}}$ how different from FT]

2.4 PropBank

PropBank (PB; **?**) is a lexicon and corpus of predicate–argument structures that takes a shallower approach than FrameNet. Whereas FrameNet frames cluster lexical predicates that evoke similar kinds of scenarios (with the same kinds of roles), and these frames are organized in a network, PropBank frames are purely lexical and there are no formal relations between different predicates or their roles. PropBank does represent lexical ambiguity—e.g., the verb *order* is ambiguous in PropBank between order-v-01 "impelled action" and order-v-02 "request to be delivered"—but PropBank's sense distinctions are generally coarser-grained than FrameNet's.

Within sense-disambiguated PropBank frames, or **rolesets**, core roles are defined with textual descriptions and assigned numbers. E.g., order-v-02 defines: **A0** "orderer", **A1** "thing ordered", **A2** "benefactive, ordered-for", and **A3** "source". Following **?**'s (**?**) theory of proto-roles, PropBank rolesets use **A0** for proto-agents and **A1** for proto-patients, but in general, there is much less consistency in interpretation of core roles across lexical predicates for PropBank than there is for FrameNet. Another difference is that PropBank's non-core roles—named **AM-***, such as **AM-PRP** for purposes— are not frame-specific.[4]

Despite all these differences, there is often a great deal in common between FrameNet-style and PropBank-style analyses, as should be apparent from comparing figure 1 and figure 2. The major benefit to PropBank is that it includes a large and comprehensively annotated corpus. We hypothesize that leveraging this large corpus indirectly can reap rewards for frame SRL.

2.5 Illinois SRL system

2.6 SemLink

[$^{\text{NS}}_{\text{S}}$ limitations!]

3 Learning from multiple domains and representations

We consider the argument identification task as statistical classification for each role of the frame in context. Let the classification input be a dependency-parsed sentence $\mathbf{x}$, the token(s) p constituting the predicate in question, and the frame f evoked by p (as determined by frame identification). Suppose we have a heuristic procedure for extracting candidate argument spans for the predicate: call this $spans(\mathbf{x}, p, f)$. We define the **local**

---

[3]Though these were *annotated* and the document level, and train/dev/test splits are by document, the frame-semantic parsing is currently restricted to the sentence level.

[4]**?** has a more extensive discussion of differences between PropBank's and FrameNet's conventions.

classifier as:

$$arg(\mathbf{x}, p, f, r) = \underset{a \in spans(\mathbf{x},p,f) \cup \{\varnothing\}}{\arg\max} \; score_{\mathbf{w}}(a \mid \mathbf{x}, p, f, r)$$

(1)

This is used to choose an argument for every role r of frame f: note that in many cases the role will be empty (non-overt), denoted $\varnothing$. The classifier chooses the best value of $score_{\mathbf{w}}$, a linear function scoring the compatibility of a candidate role–argument pair, whose parameters (feature weights) $\mathbf{w}$ are learned from data. [$^{NS}_S$ mention global inference step]

[$^{M}_K$ introduce this as a supervised domain adaptation problem? What is the source and what's the target?] [$^{M}_K$ Let D_{ft} represent the FT data and D_{ex} the exemplars?]

We experiment with several techniques for modifying the model-fitting portion of the argument identification model's local learning objective (training data, features). All of our experiments use the same form of regularization, condition on the same frame predictions[$^{NS}_S$ not oracle, right? these are SEMAFOR's current frame ID model, so not state of the art?] and syntactic preprocessing[$^{NS}_S$ does this match Dipanjan's latest experiments?], and use beam search with [$^{NS}_S$ hyperparam value] for joint decoding of arguments.[5]

[$^{NS}_S$ Domain adaptation/multitask learning techniques]

3.1 Augmenting the Training Data

3.2 Frustratingly Easy

Daumé III (2009) proposed a simple feature augmentation approach that was shown to work well in supervised domain adaptation scenarios, such as ours. We apply their method to incorporate the exemplars data by introducing a feature mapping $\Phi^{ft}(x) = (x,x)$ for the FT annotations $x \in D_{ft}$, and $\Phi^{ex}(x) = (x,0)$ for the exemplars $x \in D_{ex}$. The resulting transformed data from the two domains is then used to learn a single model. The intuition here is that by replicating the features, we allow for each feature to contribute both "general" and domain-specific weights to the model depending on whether the feature behaves similarly in both domains or not.

3.3 Guide Features

Another approach to introduce supervision for domain adaptation, which does not combine the labeled data from the two domains, to use a supervised model built on the source domain to augment features in the target domain. When training and testing in the target domain, the source domain model is effectively just a form of preprocessing to provide additional features, known as **guide features** (**??**).[6] Formally: let M_s be the model built on the source domain (for instance, the PropBank data). For every target instance x, we introduce "guide" features which use the output $M_s(x)$ obtained by applying M_s on x, which consists of the role labels assigned to various text spans in x. Two types of guide features were used, one to indicate that a span s_x was assigned a role and the other encoding the role label r_g itself. In the case where M_s produces labels that belong to the same schema as the target domain (for instance, the exemplars use the same schema as the FT annotations), we use an additional feature $f_{match}(r_t, r_g)$ to indicate that the 'guide' role label r_g of the span s_x is the same as it's true label r_t.

3.4 Type-level hierarchy features

[$^{M}_K$ NSS: notation for the frame/role/features etc?] [$^{M}_K$ How many total types of relations? Cleanup writeup based on NSS's notation.]

Frames in FrameNet are connected to each other by relations such as inheritance, temporal ordering, causality. For instance, the frame ROBBERY inherits from the more abstract frame COMMITTING_CRIME, the frame FALL_ASLEEP is preceeded by the frame BEING_AWAKE. The roles of related frames have also been mapped to indicate the correspondence between them: ROBBERY.**Perpetrator** is mapped to COMMITTING_CRIME.**Perpetrator**, which in turn maps to MISDEED.**Wrongdoer**. Frames and roles that are far apart in this hierarchy are less related than say neighbours. This hierarchy can be exploited to share information across related roles, thereby benefiting the roles that have few annotations [$^{M}_K$ say something about a greater variety of contexts is available for each role]. A simple mechanism to share information is via shared model parameters between related roles. Towards this, we define two types of hierarchical features: (1)

---

[5][$^{NS}_S$ recent work has improved upon global decoding techniques; we expect such improvements to be complementary to the gains due to the local model reported here]

[6]This is related to the technique of model stacking, where successively richer models are trained by cross-validation on the same dataset **???**.

'siblings', where roles that have a common parent share features. For every feature f_i that fires for an argument a, we add a feature which is the conjunction: $(f_i \wedge \text{parent.frame} \wedge \text{parent.role} \wedge I_{hier})$ where parent.frame = parent(frame(a)). I_{hier} is an indicator to distinguish this feature from the regular conjunction features in Semafor that use frame names and roles. (2) 'parents+siblings', where roles share features with their parents and siblings. In addition to the 'siblings' feature, we fire the following conjunction: $(f_i \wedge \text{frame} \wedge \text{frame.role} \wedge I_{hier})$. We use only the Inheritance and Subframe relations between the roles (§2.1), of which there are 4138 and 589 respectively.

We found that using more than one level of the hierarchy, does not produce any improvements in the performance at the cost of increased computation due to the greater number of features.

[$^{M}_{K}$ Write about the number of features in all models somewhere.. Baseline: 2709128, Combined training: 12972404, Frust-easy: 15681531, 'Parents+siblings': 34282958]

4 Experiments

4.1 Baseline

Our baseline system is SEMAFOR (Das et al., 2014). SEMAFOR treats the argument identification task as a structured prediction problem. It uses a linearly parametrized model that scores each candidate span given a frame element[$^{S}_{T}$ this should probably be talked about earlier?]:

$$\text{score}_{\mathbf{w}}(y \mid x) = \mathbf{w}^{\top}\mathbf{f}(x, y) \qquad (2)$$

Automatic syntactic dependency parses from MST-ParserStacked (**?**) are used to narrow the set of candidate spans considered, and as input to feature extraction.

During training, each frame element is treated as an independent log-linear classification instance, with the set of candidate spans (including the NULL span) as its output space. [$^{S}_{T}$ something about L2 regularization] But at test time, it chooses a joint assignment of all arguments that maximizes probability under the model, while satisfying the following constraints:

1. an argument may be assigned to at most one span [$^{S}_{T}$ theta criterion, Chomsky], and
2. spans of realized arguments must not overlap.

Beam search, with a beam size of 100, is used to choose the maximum joint assignment with no overlapping arguments. [$^{NS}_{S}$ does beam search require normalizing to probabilities?]

We have made several modifications to SEMAFOR's training in order to speed up experiments:

- We use squared structured hinge loss (defined below) instead of multiclass logistic regression. Using hinge loss, there is no longer a need to calculate a partition function. Gradients, and hence parameters, are sparser than in logistic regression, allowing us to use a sparse vector implementation.
- We use the online optimization method AdaDelta (Zeiler, 2012) with minibatches[$^{NS}_{S}$ minibatch size?], instead of the batch method L-BFGS (Liu and Nocedal, 1989).

We use these changes for all systems, including the baseline. While the impact on full-text performance is negligible, these changes enabled us to run more experiments with the larger exemplar dataset and expanded feature space.[7]

The details of squared hinge loss are as follows. Let $(x^{(i)}, y^{(i)})$ be the ith training example. Then the structured hinge loss on the ith example is given by:

$$Hinge_{\mathbf{w}}(i) = \\ \max_{y}\left\{\mathbf{w}^{\top}\mathbf{f}(x^{(i)}, y) + \text{cost}(y, y^{(i)})\right\} - \mathbf{w}^{\top}\mathbf{f}(x^{(i)}, y^{(i)})$$

and squared hinge loss is:

$$SquaredHinge_{\mathbf{w}}(i) = Hinge_{\mathbf{w}}(i)^2. \qquad (3)$$

In our experiments, we use $\text{cost}(y, y^{(i)}) = \mathbf{1}\{y \neq y^{(i)}\}$ [$^{S}_{T}$ define][8].

[$^{NS}_{S}$ same FN 1.5 splits as Dipanjan]

[$^{NS}_{S}$ preprocessing issues: removing duplicate sentences, merging adjacent split args in exemplars, OntoNotes PropBank preprocessing (NLTK), token-level SemLink details (such as filtering out sentences without mappable annotations; copy from WS paper)]

---

[7]With SEMAFOR's original features and training data, the result of the above changes is that full-text F_1 decreases from 59.3% to 59.1%, while training time (running optimization to convergence) decreases from 729 minutes to 82 minutes[$^{NS}_{S}$ say something about hardware this was tested on?].

[8]We experimented with recall-oriented training, where errors of omission are assigned a higher cost, but found that while recall increased, overall F_1 went down[$^{NS}_{S}$ or: failed to improve?].

4.2 Preprocessing the data

[$^{M}_{K}$ NSS: hope you're not duplicating this in the earlier section! We use the *token-level parallel annotations* from SemLink, which are available for a subset of the PB-WSJ text, hereafter referred to as SL-WSJ.] Of the available 74,977 SL-WSJ verbs, a majority cannot be mapped to FN frames for various reasons. Around 31% of the predicates have the frame label IN ("indefinite") where the mapping from VerbNet to FrameNet is ambiguous. About 20% of the instances are labeled NF ("no frame"), indicating a coverage gap in FrameNet. 21% of verbs have frame labels but no frame element annotations. Most of these are predicates with modifier arguments. Other arguments pointed to null anaphora that could not be resolved to overt arguments. This leaves 15,323 mappable instances with at least one overt argument, or 20% of SL-WSJ verbs. This is a very small subset of the entire PB annotated data.

We processed the exemplars data to remove duplicate sentences that already appear in the FT data. [$^{NS}_{S}$ tuning regularizer for all experiments] We tune the ℓ_2 regularization parameter λ on the held-out set. We searched over the following values: 10^{-5}, 10^{-7}, 10^{-9}, 10^{-12} (note that our loss is normalized). We also use the performance on the held-out set to determine the stopping criterion for the stochastic optimization.

4.3 Evaluation

The argument identification performance of the various methods is compared on two different test sets. (1) Full-text: the FrameNet 1.5 FT test split that was used in the evaluation in **?**. This data consists of sentences from 23 documents. (2) Exemplars: a randomly sampled set of sentences from the exemplars data, with approximately the same number of targets as the FT test set. Statistics of both the test sets are given in the lower half of Table 1. The FT test set has 289 unseen role types, which is much higher than the 38 in the exemplars test set. There are no unseen frame types in the exemplars test data, whereas the FT test has 46 of them. These differences are due to the manner in which the train and test splits were created, with document-level splits being used for FT and sentence-level splits for the exemplars, and also the significantly larger (hence broader coverage) exemplars training set. The last row of Table 1 shows the unseen role types faced by a model that was built on both the FT

and exemplars training data. In the FT test set it is 103, lower by ≈190 than what a FT-only model will see, thus leading us to expect that a model that combines both sets of data will certainly benefit in performance.

While the FT test set represents the benchmark set for evaluating the performance of a frame semantic role labeling system, the exemplars test set being from a different distribution of text, give us an indication of how well a model generalizes. We present precision, recall, and F1-measure microaveraged across the test instances in Table 2, for all the approaches that we tried. The first column classifies the approaches based on what resource we use and the second column indicates what training data was used. For each resource (one multi-row in the table) we show results obtained by various methods of combining it with the FT data. The $\xrightarrow{\text{guide}}$ refers to the feature augmentation discussed in Section **??**, 'Hier' represents the hierarchical features from Section **??** and 'EasyAdapt' is described in Section **??**.

To summarize the results, we find that adding the sibling-level hierarchical features, using Exemplars as gold-standard training data and incorporating PB-SRL data in the form of guide features all help in improving the performance over the baseline on both FT-test and exemplars-test.

Consulting the results in detail, adding the sibling-level hierarchical features to the baseline FT model improves the fscore by 1.2 and 1.7 points on FT-test and exemplars-test respectively, with benefits to both precision and recall. The (siblings+parents) features which consider two levels of the hierarchy produce unnoticeable benefits, suggesting that higher levels in the hierarchy can be too general [$^{M}_{K}$ give example?], resulting in very dissimilar roles firing the same hierarchy features.

SemLink, as we remarked earlier is a noisy resource and we see a confirmation for this observation in the performance of the (FT+SemLink) model, which drops the f-score by a whopping 11.2 and 16.5 points below baseline on FT-test and exemplars resp. The guide features however modulate the influence of the SemLink annotations, giving a minor increase over the baseline. [$^{M}_{K}$ resulting in a graceful but prompt exit for this resource from

---

[8]We also tried adding grand-parents and found only minor improvements. Using an expert-pruned hierarchy with relations that are most likely to help can give greater benefits, and is beyond the focus of this work.

	Full-Text	Exemplars
Sentences in training	2,780	145,838
Roles in train data	15,019	137,511
Roles (types) in training	2,644	4,821
Sentences in test data	2,420	4,132
Roles in test data	4,458	4,132
Roles (types) in test data	1,420	1,224
Unseen frames (types) in test data	46	0
Roles (types) in unseen frames	178	0
Unseen roles (types) in test data	289	38
Unseen roles (types) in test w.r.t (FT+Exemplars)	103	32

Table 1: Characteristics of the training and test data

With the exemplars resource, we find that the guide features give us a modest improvement of 1.12 in F1 on FT-test, while using it as gold-standard training data results in a bigger increase of 2.8 points. On the exemplars-test we observe a similar trend, where the (FT+Exemplars) model delivers a massive increase of ≈23 f-points and a much smaller increase is seen with the guide features. This contrasts with what we observed for SemLink. The frustratingly easy domain adaptation approach to incorporate the Exemplars further improves the F1 by a minor 0.3 on the FT-test and a decent 2.8 on the exemplars-test.

Finally, using the PB-SRL data in the form of guide features also results in small improvements over the baseline. Overall we observe that an additional resource that is very similar to the original resource helps more as training data than as a 'guide model'. Whereas the guide features help more when the additional resource is either unreliable (like SemLink) or too distinct (like PB).

[$_K^M$ Analyze the sparsity of the models? Which features have highest importance?] In addition to the performance, we also show the sizes of the various models in the column 'number of features'. All the models are quite sparse however, with ≈60-70 % of the features assigned a zero weight. The larger models are also a lot slower to train, taking about 6 times as long as the baseline.

4.4 Results

[$_S^{NS}$ results table without hierarchy features]

[$_S^{NS}$ hierarchy features: which ones work best (decided on baseline), how do they improve best result so far]

[$_S^{NS}$ Sam's curves on per-FE F_1]

[$_S^{NS}$ comparison to prior work (baseline, best result). args+frames score vs. args only]
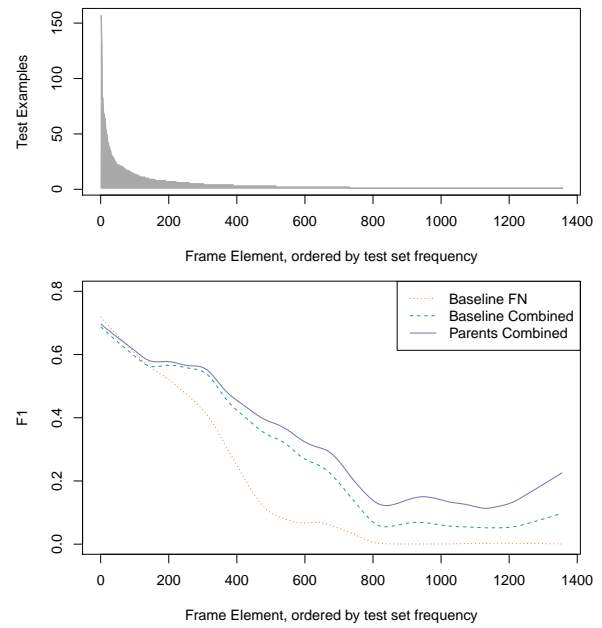

Figure 3: Count and F_1 for each frame element appearing in the test set. F_1 values have been smoothed with `loess`, with a smoothing parameter of 0.2.

[$_S^{NS}$ discussion throughout]

5 Related Work

[$_S^{NS}$ Dipanjan's other papers; mention other PB SRL work?; anything using SemLink or combining resources for SRL? see (**?**, §4)]

[$_S^{NS}$ be sure to cite: (**?**) (multiple resources), (**?**) (hierarchy)]

[$_S^{NS}$ multitask learning?]

6 Conclusion

[$_S^{NS}$ overall findings]

[$_S^{NS}$ future work: testing ground for improvements to PB (**?**) and SemLink (**?**); automatic mappings between resources]

References

Collin F. Baker, Charles J. Fillmore, and John B. Lowe. 1998. The Berkeley FrameNet project. In *Proc. of COLING-ACL*, pages 86–90. Montreal, Quebec, Canada.

Dipanjan Das, Desai Chen, André F. T. Martins, Nathan Schneider, and Noah A. Smith. 2014. Frame-semantic parsing. *Computational Linguistics*, 40(1):9–56.

Hal Daumé III. 2009. Frustratingly easy domain adaptation. *arXiv preprint arXiv:0907.1815*.

Additional Resource	Training Configuration (Features)	Number of features (approx.)	Full-Text			Exemplars		
			P	R	F_1	P	R	F_1
(Baseline)	FT (Basic)	2.7×10^6	65.57	53.82	59.12	62.63	37.65	47.03
FN Hierarchy	FT (siblings)	5.4×10^6	67.24	54.76	60.36	64.81	39.09	48.77
	FT (siblings+parents)	8.5×10^6	67.67	52.79	59.31	65.25	38.18	48.18
SemLink	SemLink $\xrightarrow{\text{guide}}$ FT	3×10^6	64.67	54.53	59.17	60.95	38.92	47.50
	FT+SemLink	5×10^6	65.50	37.80	47.90	57.15	20.80	30.50
Exemplars	Exemplars $\xrightarrow{\text{guide}}$ FT	3.5×10^6	65.24	55.96	60.24	67.71	48.08	56.23
	FT+Exemplars (Basic)	13×10^6	66.06	58.23	61.9	75.44	65.11	69.89
	FT+Exemplars (EasyAdapt)	15.7×10^6	65.70	59.04	62.19	73.88	61.40	67.06
PB-SRL	PB-SRL $\xrightarrow{\text{guide}}$ FT	3.6×10^6	64.96	54.83	59.47	61.38	39.14	47.80

Table 2: Results on two test sets: Baseline vs. individual other resources. Precision, recall, and F_1 are given as percentages.

Training Configuration (Features)	Number of features (approx.)	Full-Text			Exemplars		
		P	R	F_1	P	R	F_1
FT+Exemplars (Hier: siblings)	34.3×10^6	66.00	60.40	63.07	76.14	67.71	71.70
PB-SRL $\xrightarrow{\text{guide}}$ FT+Exemplars	17.4×10^6	67.36	58.79	62.80	77.15	65.47	70.83

Table 3: Combining best techniques across resources [$^{\text{NS}}_{\text{S}}$ TODO]

Paul Kingsbury and Martha Palmer. 2002. From Tree-Bank to PropBank. In *Proc. of LREC*, pages 1989–1993. Las Palmas, Canary Islands.

Dong C. Liu and Jorge Nocedal. 1989. On the Limited Memory BFGS Method for Large Scale Optimization. *Math. Program.*, 45(3):503–528.

Martha Palmer, Daniel Gildea, and Nianwen Xue. 2010. *Semantic Role Labeling*. Number 6 in Synthesis Lectures on Human Language Technologies. Morgan & Claypool, San Rafael, CA.

Matthew Zeiler. 2012. ADADELTA: An Adaptive Learning Rate Method. *CoRR*, abs/1212.5701. URL http://dblp.uni-trier.de/rec/bibtex/journals/corr/abs-1212-5701.